Chapter 2 : Interpretability

I Visualizat°/Analysis (of a trained network)

At the neuron level

- pick one neuron : activities on the training set → stats
  ex: classif° task

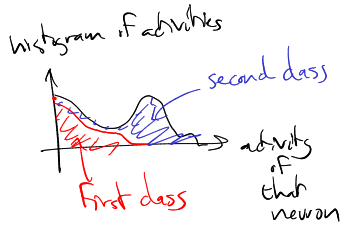

- what it sees → receptive field

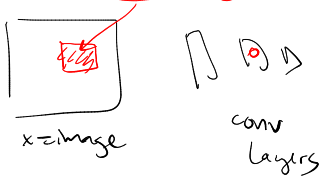

- what does it react to?
  - → display input patterns that maximize the activity
    - ↳ from the training set
  - → compute the pattern that would " " " "
    - by gradient descent: $\frac{\partial x_t}{\partial t} = \eta \frac{\partial \text{activity}(x_t)}{\partial x}$ (↑ gradient)

↳ if you apply this to the full input, looking at output neurons:
e.g. classific° task,

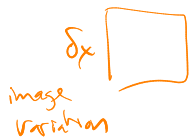

$\frac{\partial p(\text{class } c)}{\partial \text{image } x}$: sensitivity of the prediction (for class c) to the input x

$x' = x + \eta \frac{\partial p(\text{other class})}{\partial x}$ (η not much) ⇒ change completely the prediction

↳ adversarial examples [2014]
↳ adversarial attacks
↳ due to data dimension

high-dim Data space

- no dense sampling (requires $10^{\#\text{dim}}$ samples)
- all points are on the boundary
  ↳ measure concentration
  → look at uniform distrib° in the unit-ball

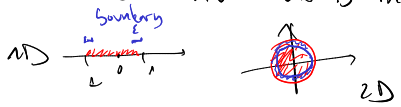

proportion of points close to the boundary → 1 (dim → ∞)

robustify

↳ train with $x \in$ training set ↦ ℓ
and $x + \delta x$ associated adversarial attack ↦ ℓ

≃ smooth function: $\frac{\partial f_{\ell,x}}{\partial x}(\delta x) = 0$

↳ $\left\| \frac{\partial f}{\partial x} \right\|^2$

adversarial attacks ⇄ robustifying techniques (loop)

- does it have an impact? $\frac{\partial f}{\partial a}(x)$ → find which neurons influence the output the most

↑ activity of that neuron

⇒ reasoning at the neuron level: compare to human brain

The case of CNN

- which parts of the image are responsible for the decision?

grad-CAM: Class Activation Maps → classif° pb

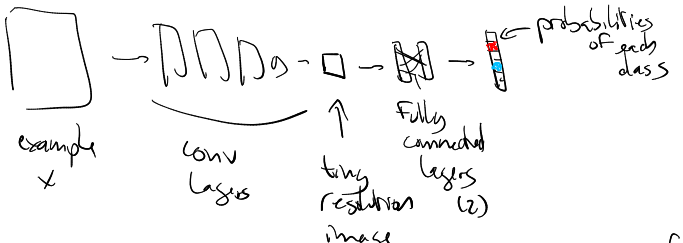

- pick a class c

- activities of that layer: A^k_{ij} ← feature, ← pixel location

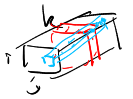

$\mapsto p(c) = y_c$

- importance of feature k for class c:

$$\alpha^c_k = \frac{1}{\#\text{pixels}} \sum_{ij} \frac{\partial y_c}{\partial A^k_{ij}} \in \mathbb{R}$$

- importance of one pixel (i,j):

$$\sum_k \alpha^c_k A^k_{ij} \in \mathbb{R}$$

$$- \text{ReLU}\left(\sum_k \alpha^c_k A^k_{ij}\right)$$

⇓

heat map for class c

← on the tiny resolution ← Scale it up to the original resolution ← Blurry